

1) OpenMP Parallel Sparse Matrix Vector Multiplication

You might notice that I do not set the number of threads in the code. I set the number of threads by requesting the according number of cores in the job script, which I dynamically generate from a template. I wrote the parallel code in a general way, letting the OpenMP runtime choose the amount of threads it spawns. Usually for the code I wrote that should be equal to the number of available cores.

a) Parallel Matrix Conversion Code

I found it to be quite hard to achieve a significant speedup here. Most loops are not worth paralyzing. The most significant speedup arises from paralyzing the loop which calculates the number of non zero elements per block for the CSB format. Here I simply paralyze the outer loop which iterates over all columns.

The loop after that which allocates the memory for all blocks is also worth paralyzing. To make things a little more efficient it makes sense to collapse the two most outer loops. That leads to each individual block being a chunk of work that can be executed by a single thread.

I also paralyzied the following loop but it doesn't seem to make a significant difference.

b) Parallel Sparse Matrix Vector Multiplication using the CSC Format

Simply paralyzing the outer loop iterating over all columns leads to the best results of all different ways I have tested. It seems like for large matrices there are enough columns to efficiently distribute the work and it keeps the overhead very low since there are no data conflicts when writing to the b output vector in this way. However, since the time measurements are biased due to cache utilization of the parallel version after executing the sequential version on the same input, it is hard to provide a solid assessment here.

c) Parallel Sparse Matrix Vector Multiplication using the CSB Format

This one was actually the most complicated to efficiently paralyze. The most efficient way I found is based on paralyzing all blocks by collapsing the two outer loops. However, this creates writing conflicts into the output vector between all threads. The most efficient way to resolve this is to provide each thread with it's own copy of an empty output vector and summing them up afterwards, which is exactly what I did.

d) Results for Three Large Matrices

The requested plots are provided by Figure 1 for the *it-2004* matrix, Figure 2 for *sk-2005* and Figure 3 for the *twitter7* matrix. All plots contain a horizontal black line as a reference for no speedup (y -value of 1).




Figure 1: Results for the *it-2004* matrix.




Figure 2: Results for the *sk-2005* matrix.

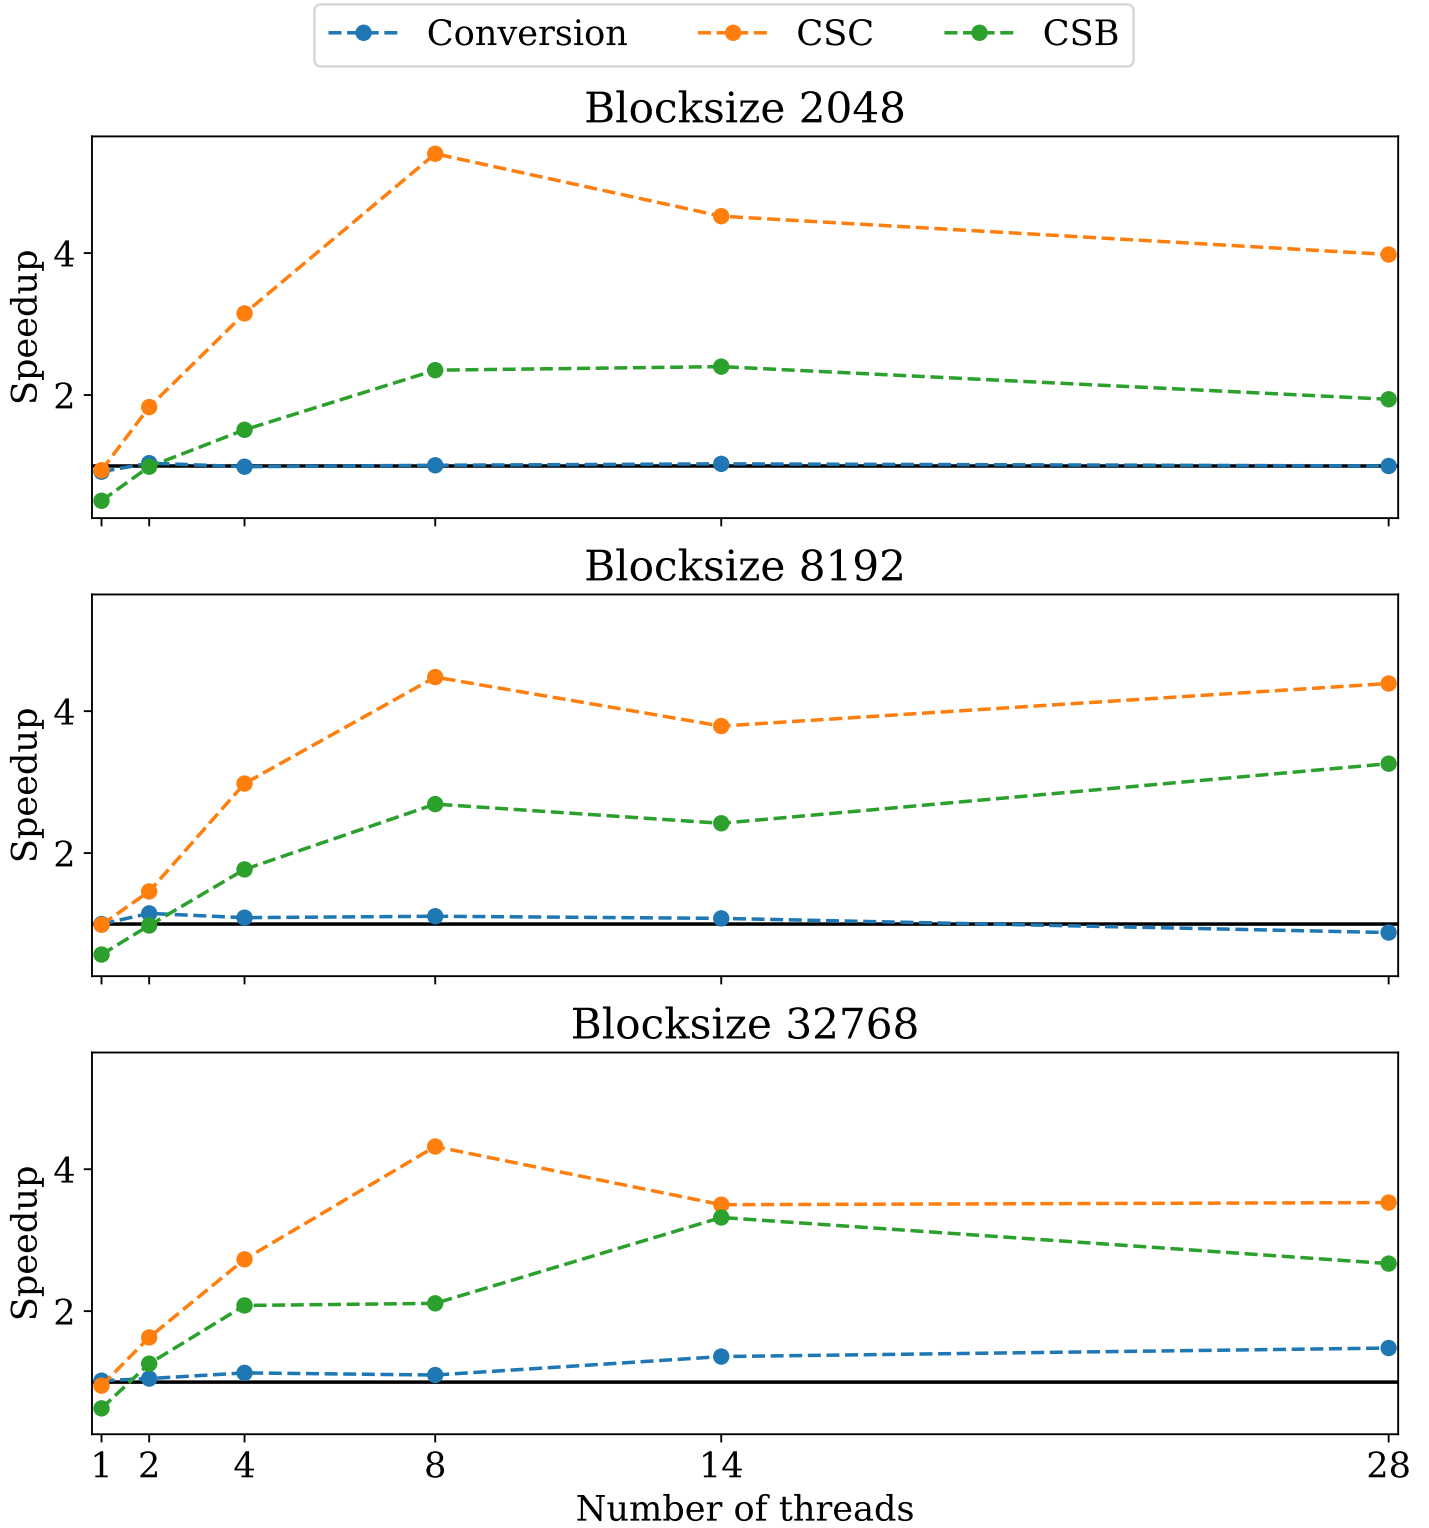


Figure 3: Results for the *twitter7* matrix.

Over all the results are inconclusive and can't fully be trusted because of caching effects. If I understand the methods correctly, the CSC multiplication results should not depend on the blocking factor at all. However, they do because the CSB sequential routine is executed in between of the CSC sequential and parallel methods, effecting the caching.

Nevertheless, there is a general trend we can observe: All methods do not scale well beyond 14 threads. This can simply be explained by the overhead becoming to expansive when adding more and more threads. There simply isn't enough work per thread to justify the amount of overhead generated by each thread when using a high number of threads.